

# Trees, Bagging, and Boosting

He He  
(adapted from David Rosenberg's slides)

CDS, NYU

April 14, 2019


# Contents

- 1 Decision Trees
  - Input Space Partition
  - Node Splitting
  - Trees in General
- 2 Bagging and Random Forests
  - Variance of an Estimator
  - The Benefits of Averaging
  - Bootstrap
  - Bagging
  - Random Forests
- 3 Boosting
  - Adaboost: The Algorithm

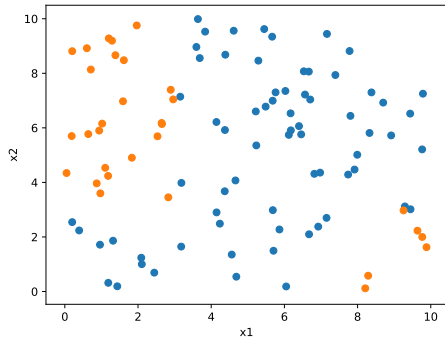
# Decision Trees

## Motivating example in 2d

- Partition data into different (axis-aligned) regions **recursively**



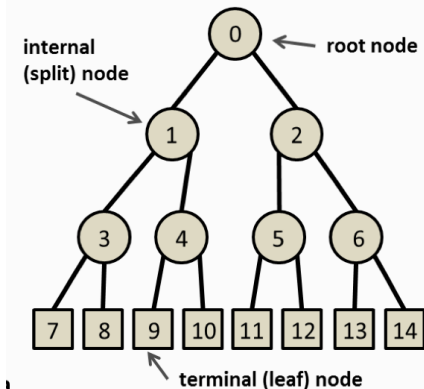
# Classification flowchart



Is this a linear or non-linear classifier?

# Decision trees setup

## A general tree structure



We'll only consider

- **binary** trees (vs multiway trees where nodes can have more than 2 children)
- each node contains a subset of data points
- decisions at each node involve only a **single** feature (i.e. input coordinate)
- for continuous variables, splits always of the form

$$x_i \leq t$$

- for discrete variables, partitions values into two groups

# Regularization of decision trees

- What will happen if we keep splitting the data?
  - Every data point will be in its own region—**overfitting**.
- When to stop splitting? (control complexity of the hypothesis space)
  - Limit number of total nodes.
  - Limit number of terminal nodes.
  - Limit tree depth.
  - Require minimum number of data points in a terminal node.
  - **Backward pruning** – the approach of **CART** (Breiman et al 1984):
    - ① Build a really big tree (e.g. until all regions have  $\leq 5$  points).
    - ② **Prune** the tree back greedily all the way to the root, assessing performance on validation.

# How to split

**Goal** Find a tree that minimize the task loss (e.g., squared loss) within a given complexity.

**Problem** Finding the optimal binary is computationally intractable.

**Solution** Greedy algorithm.

- Find the best split (according to some criteria) for a terminal node (initially the root)
- Add two children nodes
- Repeat until a stopping criterion is reached (e.g., max depth)



# Evaluate splits

Let's think about what makes a good split.

Which one is better?

Split 1  $R_1 : 8+ / 2- \quad R_2 : 2+ / 8-$

Split 2  $R_1 : 6+ / 4- \quad R_2 : 1+ / 9-$

Which one is better?

Split 1  $R_1 : 8+ / 2- \quad R_2 : 2+ / 8-$

Split 2  $R_1 : 6+ / 4- \quad R_2 : 0+ / 10-$

In general, we want to produce **pure** nodes, i.e. close to single-class node.

# Misclassification error in a node

Let's formalize things a bit.

- Consider classification case:  $\mathcal{Y} = \{1, 2, \dots, K\}$ .
- What's in a node?
  - Let node  $m$  represent region  $R_m$ , with  $N_m$  observations
  - Denote proportion of observations in  $R_m$  with class  $k$  by

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{\{i: x_i \in R_m\}} 1(y_i = k).$$

- Predict the majority class in node  $m$ :

$$k(m) = \arg \max_k \hat{p}_{mk}.$$

- Misclassification rate in node  $m$ :

$$1 - \hat{p}_{mk(m)}.$$

# Node Impurity Measures

How to quantify impurity?

- Three measures of **node impurity** for leaf node  $m$ :

Misclassification error

$$1 - \hat{p}_{mk(m)}.$$

Gini index

$$\sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

Entropy

$$-\sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}.$$

- Gini index and entropy work well in practice.

## Impurity of a split

A potential split produces two nodes,  $R_L$  and  $R_R$ . How do we score it?

- Suppose we have  $N_L$  points in  $R_L$  and  $N_R$  points in  $R_R$ .
- Let  $Q(R_L)$  and  $Q(R_R)$  be the node impurity measures for each node.
- Then find split that minimizes the **weighted average of node impurities**:

$$\frac{N_L Q(R_L) + N_R Q(R_R)}{N_L + N_R}$$


Example:

$R_1 : 8 + /2 -$      $R_2 : 1 + /4 -$

What's the weighted misclassification rate?

## Two-Class Node Impurity Measures

Consider binary classification. Let  $p$  be the relative frequency of class 1.



Misclassification error is not strictly concave thus may not guarantee improvement over the parent node.

# Finding the Split Point

How to find a split point that minimizes a given impurity measure?

- Consider splitting on the  $j$ 'th feature  $x_j$ .
- If  $x_{j(1)}, \dots, x_{j(n)}$  are the sorted values of the  $j$ 'th feature,
  - we only need to check split points between adjacent values
  - traditionally take split points halfway between adjacent values:

$$s_j \in \left\{ \frac{1}{2} (x_{j(r)} + x_{j(r+1)}) \mid r = 1, \dots, n-1 \right\}. \quad n-1 \text{ splits} \quad (1)$$

- Enumerate  $d$  features and  $n-1$  split points for each feature.

# Regression trees

- Predict the mean value of a node

$$k(m) = \text{mean}(y_i \mid x_i \in R_m). \quad (2)$$

- Squared loss as the node impurity measure.
- Everything else remains the same as classification trees.

# Categorical features


- For a categorical feature, we split its values into two groups.
- Given a set of categories of size  $k$ , how many distinct splits? (its power set)
- Finding the optimal split is **intractable** in general.
- Approximations
  - Numeric encoding Random or proportion of class 0 (for binary classification)
  - One-hot encoding May grow imbalanced trees, e.g., left-branching
  - Binary encoding Robust to large cardinality
- Statistical issues with categorical features
  - If a category has a very large number of categories, we can **overfit**.
  - Extreme example: Row Number could lead to perfect classification with a single split.



- Trees are certainly easier to explain than other classifiers.
- Can be used to discover non-linear features.
- Small trees seem interpretable. For large trees, maybe not so easy.
- Approximate neural network decision boundaries to gain interpretability
  - Wu M, Hughes M, Parbhoo S, Zazzi M, Roth V, Doshi-Velez F. [Beyond Sparsity: Tree Regularization of Deep Models for Interpretability](#). Association for the Advancement of Artificial Intelligence (AAAI). 2018

# Trees vs linear models

Trees have to work much harder to capture linear relations.



## Decision trees:

- **Non-linear** classifier that recursively partitions the input space.
- Non-metric: make no use of geometry, i.e. no inner-product or distances.
- Non-parametric: make no assumption of the data distribution.

## Pros:

- Simple to understand.
- Interpretable, feature selection for free.

## Cons:

- Poor linear modeling.
- Unstable / high variance, tend to **overfit**. → Next, how to fix this.

## Bagging and Random Forests

## Recap: statistic and point estimator

- Observe data  $\mathcal{D} = (x_1, x_2, \dots, x_n)$  sampled i.i.d. from a parametric distribution  $p(\cdot | \theta)$ .
- A **statistic**  $s = s(\mathcal{D})$  is any function of the data.
  - E.g., sample mean, sample variance, histogram, empirical data distribution
- A statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  is a **point estimator** of  $\theta$  if  $\hat{\theta} \approx \theta$ .

### Review questions

In frequentist statistics,

- Is  $\theta$  random?
- Is  $\hat{\theta}$  random?
- Is the function  $s(\cdot)$  random?

## Recap: bias and variance of an estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- What are some parameters of the sampling distribution we might be interested in?

**Bias**  $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}] - \theta.$

**Variance**  $\text{Var}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}^2] - \mathbb{E}^2[\hat{\theta}].$

- Is bias and variance random?
  - Neither bias nor variance depend on a specific sample  $\mathcal{D}_n$ . We are **taking expectation over  $\mathcal{D}$** .
- Why do we care about variance?
  - $\hat{\theta}(\mathcal{D}) = x_1$  is an unbiased estimator of the mean of a Gaussian, but would be farther away from  $\theta$  than the sample mean.

## Variance of a Mean

Using a single estimate may have large standard error

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator:  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- We could use a single estimate  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .
- The standard error is  $\sqrt{\text{Var}(\hat{\theta})} = \sigma$ .

Average of estimates has smaller standard error

- Consider a new estimator that takes the average of i.i.d.  $\hat{\theta}_1, \dots, \hat{\theta}_n$  where  $\hat{\theta}_i = \hat{\theta}(\mathcal{D}^i)$ .
- Average has the same expected value but smaller standard error:

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \theta \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n} \quad (3)$$

# Averaging Independent Prediction Functions

Let's apply **averaging** to reduce variance of prediction functions.

- Suppose we have  $B$  independent training sets from the same distribution ( $\mathcal{D} \sim p(\cdot | \theta)$ ).
- Learning algorithm (estimator) gives  $B$  decision functions:  $\hat{f}_1(x), \hat{f}_2(x), \dots, \hat{f}_B(x)$
- Define the average prediction function as:

$$\hat{f}_{\text{avg}} \stackrel{\text{def}}{=} \frac{1}{B} \sum_{b=1}^B \hat{f}_b \quad (4)$$

- What's random here?
- The  $B$  independent training sets are random, which gives rise to variation among the  $\hat{f}_b$ 's.
- **Concept check:** What's the distribution of  $\hat{f}$  called? What do we know about the distribution?



## Averaging reduce variance of predictions

- The average prediction on  $x_0$  is

$$\hat{f}_{\text{avg}}(x_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x_0).$$

- $\hat{f}_{\text{avg}}(x_0)$  and  $\hat{f}_b(x_0)$  have the same expected value, but
- $\hat{f}_{\text{avg}}(x_0)$  has smaller variance (see 3):

$$\text{Var}(\hat{f}_{\text{avg}}(x_0)) = \frac{1}{B} \text{Var}(\hat{f}_1(x_0))$$

- **Problem:** in practice we don't have  $B$  independent training sets...

# The Bootstrap Sample

How do we simulate multiple samples when we only have one?

- A **bootstrap sample** from  $\mathcal{D}_n = (x_1, \dots, x_n)$  is a sample of size  $n$  drawn **with replacement** from  $\mathcal{D}_n$ .
- Some elements of  $\mathcal{D}_n$  will show up multiple times, and some won't show up at all.

How similar are the bootstrap samples?

- Each  $x_i$  has a probability of  $(1 - 1/n)^n$  of not being selected.
- Recall from analysis that for large  $n$ ,

$$\left(1 - \frac{1}{n}\right)^n \approx \frac{1}{e} \approx .368.$$

(5)

- So we expect  $\sim 63.2\%$  of elements of  $\mathcal{D}_n$  will show up at least once.

# The Bootstrap Method

## Definition

A **bootstrap method** is when you **simulate** having  $B$  independent samples from  $P$  by taking  $B$  bootstrap samples from the sample  $\mathcal{D}_n$ .

- Given original data  $\mathcal{D}_n$ , compute  $B$  bootstrap samples  $D_n^1, \dots, D_n^B$ .
- For each bootstrap sample, compute some function

$$\phi(D_n^1), \dots, \phi(D_n^B)$$

- Work with these values as though  $D_n^1, \dots, D_n^B$  were i.i.d. samples from  $P$ .
- **Amazing fact:** This is often very close to what we'd get with independent samples from  $P$ .

# Independent vs Bootstrap Samples

- Want to estimate  $\alpha = \alpha(P)$  for some unknown  $P$  and some complicated  $\alpha$ .
- Point estimator  $\hat{\alpha} = \hat{\alpha}(\mathcal{D}_{100})$  for samples of size 100.
- Histogram of  $\hat{\alpha}$  based on
  - 1000 independent samples of size 100, vs
  - 1000 bootstrap samples of size 100




Figure 5.10 from [ISLR](#) (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

## Side note: Bootstrap in Practice

We can use bootstrap to get error bars in a cheap way.

- Suppose we have an estimator  $\hat{\theta} = \hat{\theta}(\mathcal{D}_n)$ , e.g., the accuracy of your classifier.
- To get error bars, we can compute the “bootstrap variance”.
  - Draw  $B$  bootstrap samples.
  - Compute sample variance of  $\hat{\theta}(\mathcal{D}_n^1), \dots, \hat{\theta}(\mathcal{D}_n^B)$ ..
  - Could report

$$\hat{\theta}(\mathcal{D}_n) \pm \sqrt{\text{Bootstrap Variance}}$$

# Ensemble methods

## Key ideas:

- **Averaging** i.i.d. estimates reduces variance without making bias worse.
- Can use bootstrap to simulate multiple data samples.

## Ensemble methods:

- Combine outputs from multiple models.
  - Same learner on different datasets: ensemble + bootstrap = bagging.
  - Different learners on one dataset: they may make similar errors.
- Parallel ensemble: models are built independently, e.g., bagging
- Sequential ensemble: models are built sequentially, e.g., boosting
  - Try to add new learners that do well where previous learners lack

- Draw  $B$  bootstrap samples  $D^1, \dots, D^B$  from original data  $\mathcal{D}$ .
- Let  $\hat{f}_1, \hat{f}_2, \dots, \hat{f}_B$  be the prediction functions from training on  $D^1, \dots, D^B$ , respectively.
- The **bagged prediction function** is a **combination** of these:

$$\hat{f}_{\text{avg}}(x) = \text{Combine} \left( \hat{f}_1(x), \hat{f}_2(x), \dots, \hat{f}_B(x) \right)$$

- How might we combine
  - prediction functions for regression?
  - binary class predictions?
  - binary probability predictions?
  - multiclass predictions?

# Out-of-Bag Error Estimation

- Each bagged predictor is trained on about 63% of the data.
- Remaining 37% are called **out-of-bag (OOB)** observations.
- For  $i$ th training point, let

$$S_i = \{b \mid D^b \text{ does not contain } i\text{th point}\}.$$

- The **OOB prediction** on  $x_i$  is

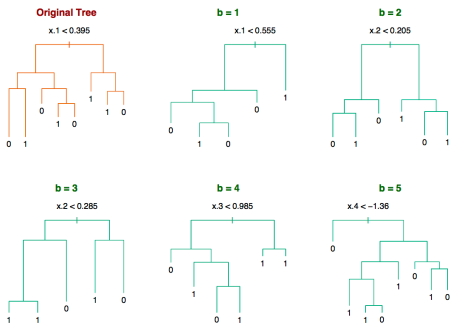
$$\hat{f}_{\text{OOB}}(x_i) = \frac{1}{|S_i|} \sum_{b \in S_i} \hat{f}_b(x_i).$$

- The OOB error is a good estimate of the test error.
- OOB error is similar to cross validation error – both are computed on training set.



# Bagging Classification Trees

- Input space  $\mathcal{X} = \mathbf{R}^5$  and output space  $\mathcal{Y} = \{-1, 1\}$ . Sample size  $n = 30$ .



- Each bootstrap tree is quite different: different splitting variable at the root
- **High variance:** high degree of model variability from small perturbations of the training data.
- Conventional wisdom: Bagging helps most when base learners are relatively unbiased but has high variance / low stability  $\implies$  decision trees.

From HTF Figure 8.9

## Variance of a Mean of Correlated Variables

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  i.i.d. with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- Suppose  $\forall i \neq j, \text{Corr}(\hat{\theta}_i, \hat{\theta}_j) = \rho$ . Then

$$\text{Var}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \rho\sigma^2 + \frac{1-\rho}{n}\sigma^2.$$

- For large  $n$ , the  $\rho\sigma^2$  term dominates – limits benefit of averaging.

# Correlation between bootstrap samples


- Averaging  $\hat{f}_1, \dots, \hat{f}_B$  reduces variance if they're based on **i.i.d.** samples from  $P_{\mathcal{X} \times \mathcal{Y}}$
- Bootstrap samples are
  - independent samples from the training set, but
  - are **not** independent samples from  $P_{\mathcal{X} \times \mathcal{Y}}$ .
- This dependence limits the amount of variance reduction we can get.
- Solution: reduce the dependence between  $\hat{f}_i$ 's.

## Key idea

Use bagged decision trees, but modify the tree-growing procedure to reduce the dependence between trees.

- Build a collection of trees independently (in parallel).
- When constructing each tree node, restrict choice of splitting variable to a randomly chosen subset of features of size  $m$ .
  - Avoid dominance by strong features.
- Typically choose  $m \approx \sqrt{p}$ , where  $p$  is the number of features.
- Can choose  $m$  using cross validation.

# Random Forest: Effect of $m$ size



From [An Introduction to Statistical Learning, with applications in R](#) (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

- Usual approach is to build very deep trees—low bias but **high variance**
- Ensembling many models reduces variance
  - Motivation: Mean of i.i.d. estimates has smaller variance than single estimate.
- Use bootstrap to simulate many data samples from one dataset
  - $\implies$  Bagged decision trees
- But bootstrap samples (and the induced models) are correlated.
- Bagging seems to work better when we are combining a diverse set of prediction functions.
  - $\implies$  random forests (randomized tree building)

# Boosting

**Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel.

**Boosting** Reduce the error rate of a high bias estimator by ensembling many estimators trained in sequential.

- A **weak/base learner** is a classifier that does slightly better than random.
- Weak learners are like “rules of thumb”:
  - “Viagra”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different examples (**reweighted data**)
  - Weak learners have different contributions to the final prediction (**reweighted classifier**)



# AdaBoost: Setting

- **Binary** classification:  $\mathcal{Y} = \{-1, 1\}$
- Base hypothesis space  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{-1, 1\}\}$ .
- Typical base hypothesis spaces:
  - **Decision stumps** (tree with a single split)
  - Trees with few terminal nodes
  - Linear decision functions

# Weighted Training Set

Each base learner is trained on weighted data.

- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Weights  $(w_1, \dots, w_n)$  associated with each example.
- **Weighted empirical risk:**


$$\hat{R}_n^w(f) = \frac{1}{W} \sum_{i=1}^n w_i \ell(f(x_i), y_i) \quad \text{where } W = \sum_{i=1}^n w_i$$

- Examples with larger weights have more influence on the loss.

# AdaBoost - Rough Sketch

- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Start with equal weight on all training points  $w_1 = \dots = w_n = 1$ .
- Repeat for  $m = 1, \dots, M$ :
  - Find base classifier  $G_m(x)$  that tries to fit weighted training data (but may not do that well)
  - Increase weight on the points  $G_m(x)$  misclassifies
- So far, we've generated  $M$  classifiers:  $G_1, \dots, G_M : \mathcal{X} \rightarrow \{-1, 1\}$ .

# AdaBoost: Schematic



From ESL Figure 10.1

# AdaBoost - Rough Sketch

- Training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .
- Start with equal weight on all training points  $w_1 = \dots = w_n = 1$ .
- Repeat for  $m = 1, \dots, M$ :
  - Base learner fits weighted training data and returns  $G_m(x)$
  - Increase weight on the points  $G_m(x)$  misclassifies
- Final prediction  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ . (recall  $G_m(x) \in \{-1, 1\}$ )
- What are desirable  $\alpha_m$ 's?
  - nonnegative
  - larger when  $G_m$  fits its weighted  $\mathcal{D}$  well
  - smaller when  $G_m$  fits weighted  $\mathcal{D}$  less well

# Adaboost: Weighted Classification Error

- Weights of base learners depend on their performance. How to evaluate each base learner?
- In round  $m$ , base learner gets a weighted training set.
  - Returns a base classifier  $G_m(x)$  that minimizes weighted 0–1 error.
- The **weighted 0-1 error** of  $G_m(x)$  is

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i 1(y_i \neq G_m(x_i)) \quad \text{where } W = \sum_{i=1}^n w_i.$$

- Notice:  $\text{err}_m \in [0, 1]$ .

# AdaBoost: Classifier Weights

- The weight of classifier  $G_m(x)$  is  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$ .



- Higher weighted error  $\implies$  lower weight

## Adaboost: Example Reweighting

We want the base learner to focus more on examples misclassified by the previous learner.

- We train  $G_m$  to minimize weighted error, and it achieves  $\text{err}_m$ .
- Then  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$  is the weight of  $G_m$  in final ensemble.
- Suppose  $w_i$  is weight of example  $i$  before training:
  - If  $G_m$  classifies  $x_i$  correctly, then  $w_i$  is unchanged.
  - Otherwise,  $w_i$  is increased as


$$\begin{aligned} w_i &\leftarrow w_i e^{\alpha_m} \\ &= w_i \left( \frac{1 - \text{err}_m}{\text{err}_m} \right) \end{aligned}$$

- For  $\text{err}_m < 0.5$  (weak learner), this always increases the weight.



## Adaboost: Example Reweighting

- Any misclassified point has weight adjusted as  $w_i \leftarrow w_i \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .



- The smaller  $\text{err}_m$ , the more we increase weight of misclassified points.

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute **weighted empirical 0-1 risk**:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq G_m(x_i)) \quad \text{where } W = \sum_{i=1}^n w_i.$$

- ③ Compute **classifier weight**:  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .
  - ④ Update **example weight**:  $w_i \leftarrow w_i \cdot \exp[\alpha_m \mathbf{1}(y_i \neq G_m(x_i))]$
- ③ Return **voted classifier**:  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

# AdaBoost with Decision Stumps

- After 1 round:




Figure: Plus size represents weight. Blackness represents score for red class.

# AdaBoost with Decision Stumps

- After 3 rounds:




Figure: Plus size represents weight. Blackness represents score for red class.

# AdaBoost with Decision Stumps

- After 120 rounds:





Figure: Plus size represents weight. Blackness represents score for red class.

# Typical Train / Test Learning Curves


- Might expect too many rounds of boosting to overfit:



From Rob Schapire's NIPS 2007 Boosting tutorial.

# Learning Curves for AdaBoost

- In typical performance, AdaBoost is surprisingly resistant to overfitting.
- Test continues to improve even after training error is zero!



# Summary

- Shallow decision tree + boosting
  - “best off-the-shelf classifier in the world” (Hastie et al., 2009)
  - Used in the first successful real-time face detector (Viola and Jones, 2001)
  - **XGBoost**: very popular in competitions
- Next week
  - What is the objective function of Adaboost?
  - Generalize to other loss functions.